

Day 3 lab: from a trained model to the microcontroller

Naeemullah Khan

naeemullah.khan@kaust.edu.sa



جامعة الملك عبد الله
للعلوم والتقنية
King Abdullah University of
Science and Technology

KAUST Academy
King Abdullah University of Science and Technology

October 5, 2026


- ▶ **Goal:** your group has a motion classifier that runs on the XIAOML Kit, and you know its flash use, its arena size, and its latency.
- ▶ **Board:** the XIAOML Kit, with the Arduino firmware.
- ▶ **Deliverable:** a live demonstration of the four motions, and the file `report.md` with the comparison table of the two paths and the Decision Log.

From the lecture to the lab

The lecture gave the format, the runtime, and the features. In the lab, you train a model on your Day 2 dataset, you convert it, you write the inference code, and you measure the three numbers on the board.

Key point: A number from the laptop is an estimate. Only the number from the board is a result.

Plan of the three hours



Part	Work	Time
A	Train: features of the Day 2 dataset, a small classifier	45 min
B	Convert and deploy: LiteRT file, C array, sketch with TensorFlow Lite Micro	50 min
C	Measure: flash, arena, latency, with PSRAM and with no PSRAM	30 min
D	Compare: the same task in Edge Impulse Studio	25 min
Check by the instructor		30 min

Write each result in `report.md` when you get it.

Hardware for each group

- ▶ XIAOML Kit with the Arduino firmware. If MicroPython is on the board, do the last step of the Day 2 lab first.
- ▶ USB-C data cable.
- ▶ Laptop with your Day 2 dataset in `Labs/day02/data/`.


Software

- ▶ Arduino IDE 2 with the esp32 core 3.3.12.
- ▶ New today: the library `Chirale_TensorFlowLite 2.0.0`.
- ▶ New today: the Python packages `tensorflow-cpu` and `ai-edge-litert`.
- ▶ Your Edge Impulse project of Day 2.

Lab files `Labs/day03/`. Run the commands from this folder.

Install In the environment of Day 1: `pip install tensorflow-cpu ai-edge-litert`

Two paths to the same classifier




- ▶ The two paths use the same recordings, the same window of 2 s, and the same type of features: 63 values.
- ▶ In the first path, you write the inference code, and you see each step.
- ▶ In the second path, a cloud tool makes the features, the model, and the library.
- ▶ At the end, one table compares the two results on the board.

Part A: train (45 min)

Start the notebook: `../../.venv/bin/jupyter lab motion_classifier.ipynb`. It uses your recordings in `Labs/day02/data/`.

Task	You write	The notebook checks
1	The function that removes the mean and gives the RMS	A sine wave with the amplitude 1 gives an RMS of 0.707
2	The two hidden layers of the model	The model has 1534 parameters
3	Your prediction: features against the raw window	Not checked. You compare it with the result.
4	The two lines of the converter (Part B)	The file starts correctly, and its output is equal to the output of Keras
5	Your estimate of the arena size (Part B)	Not checked. You measure it in Part C.



- ▶ **You see:** the plot of the spectral power for one window of each class. Find the axis and the frequency that separate your classes.
- ▶ **You record:** the training accuracy, the test accuracy, and the confusion matrix.

Prediction (task 3). The notebook trains a second model with the raw window: 300 inputs. How many parameters does it have? Which input gives the higher test accuracy? Write the prediction before you run section 6.

Task 4. Write the two lines of the converter in section 7 of the notebook. The notebook then prints lines of this form:

```
Task 4: complete
LiteRT file:          ... bytes
operator types:       [...]
largest difference between Keras and LiteRT: ...
```

- ▶ The difference must be smaller than 0.00001.
- ▶ **You record:** the file size and the operator types. Task B1 needs the operator types.

Section 8 writes three files into `sketches/motion_classifier/`:

<code>model.h</code>	The LiteRT file as a C array
<code>model_settings.h</code>	The class names, and the mean and the standard deviation of each feature
<code>test_window.h</code>	One test window, its 63 features, and the output of the model for it

Task 5. Estimate the arena size from the table of the last cell. Write the estimate in the report before Part C.

Install the library Chirale_TensorFlowLite 2.0.0. Open `sketches/motion_classifier/motion_classifier.ino`. Set Tools > PSRAM > Disabled.

Task	Place	You write
B1	<code>setupModel()</code>	The operator resolver: one <code>Add...()</code> call for each operator type of your model
B2	<code>setupModel()</code>	The interpreter, and the call of <code>AllocateTensors()</code>
B3	<code>classify()</code>	Normalize the features, write the input tensor, call <code>Invoke()</code> , read the output tensor
B4	<code>bestClass()</code>	The position of the largest probability

- ▶ The sketch compiles before you start. It then prints Tasks B1 and B2 are not complete.
- ▶ **You record:** the line `Sketch` uses ... bytes of this first build. It is the flash use with no runtime in the program.

Upload the complete sketch. The Serial Monitor at 115200 baud shows:

```
Arena location:      internal RAM
Model size (flash):  ... bytes
Sketch size (flash): ... bytes
Arena size:          4096 bytes
Arena used:          ... bytes
Self-test, features: largest difference ...
Self-test, model:    largest difference ...
Self-test, class:    ... (expected: ...)
Self-test:           PASS
```

- ▶ The self-test calculates the features of the window of `test_window.h` on the board, runs the model, and compares the results with the results of the notebook.
- ▶ FAIL: check task B3, and run section 8 of the notebook again.
- ▶ Then make each of the four motions. The display shows the class and its probability.
- ▶ **You record:** for each motion, the class that the board shows most of the time.

Source: The output shows the form of the lines.

Part C: measure three numbers (30 min)

Number	Where you read it	You record
Flash	Sketch uses ... bytes of the build, and Model size of the Serial Monitor	The two values, and the difference to the first build of Part B
Arena	Arena used of the Serial Monitor	The value, and your estimate of task 5
Latency	The columns features_us and invoke_us	The median and the maximum of 20 lines

One line of the Serial Monitor for each inference:

```
class,probability,features_us,invoke_us,late_samples  
.....
```

- ▶ The sum of the two times must be below 20 000 microseconds: one sample period.
- ▶ late_samples counts the samples that the loop read too late. Write the value after 20 lines.

The arena size

1. Set `kTensorArenaSize` to 1024. Upload. Write the message of the Serial Monitor in the report.
2. Set it to the value `Arena` used plus 256 bytes. Upload. The self-test must pass.
3. Compare the measured value with your estimate. Explain the difference in one sentence.

The second configuration

1. Set `Tools > PSRAM > OPI PSRAM`.
2. Change the line of the sketch to `#define ARENA_IN_PSRAM 1`.
3. Upload. The Serial Monitor prints `Arena` location: `PSRAM`.
4. Write the same three numbers in the second column of the report.

Prediction. Is the inference faster, slower, or equal when the arena is in the PSRAM? Write your answer and the reason before you measure.

1. **Impulse.** Open your project of Day 2. In Create impulse: window size 2000 ms, window increase 200 ms, the block Spectral Analysis, and the block Classification.
 2. **Features.** In Spectral features: FFT length 32. Generate the features.
 3. **Training.** In Classifier: two dense layers with 20 and 10 neurons, 30 training cycles. Then run Model testing.
 4. **Library.** In Deployment: Arduino library. Add the ZIP file with Sketch > Include Library > Add .ZIP Library.
 5. **Sketch.** Open sketches/ei_motion_inference/. Change the first `#include` line to the header of your library. Upload.
- ▶ **You record:** the test accuracy, and the estimates of the Studio for the inferencing time, the peak RAM, and the flash use.
 - ▶ **You record:** the two times of the line Predictions (DSP: ... ms, Classification: ... ms, ...) and the line Sketch uses ... bytes.
 - ▶ If the build of the library is not complete in time, use the estimates of the Studio.

Number	Your sketch with TensorFlow Lite Micro	Library of Edge Impulse
Test accuracy	Section 5 of the notebook	Model testing
Flash	Sketch uses ... bytes	Sketch uses ... bytes, or the estimate
RAM	Arena used	The peak RAM of the Studio
Time for the features	features_us	DSP time
Time for the inference	invoke_us	Classification time
Kernels	Reference kernels	Optimized kernels (ESP-NN)

- ▶ Write “measured” or “estimate” next to each number of the second column.
- ▶ The two paths use different units: microseconds and milliseconds. Write the unit.

The instructor checks this at the end of the lab:

- ▶ ✓ The notebook prints `complete` for tasks 1, 2, and 4.
- ▶ ✓ The Serial Monitor shows `Self-test: PASS`.
- ▶ ✓ The board shows the correct class for each of the four motions.
- ▶ ✓ The report has the three predictions: task 3, the arena estimate, and the PSRAM question. You wrote them before the measurement.
- ▶ ✓ The comparison table has measured numbers for your sketch: flash, arena, and latency, with PSRAM and with no PSRAM.
- ▶ ✓ The comparison table has numbers for the path of Edge Impulse. Each number says if it is a measurement or an estimate.
- ▶ ✓ The Decision Log gives numbers and names one trade-off.

Hand in: the file `report.md`. Show the kit to the instructor.

Write about 100 words. State one design decision, give your measured numbers, and name the trade-off.

A product must classify the motion of a package for one year with a battery. Which path do you select for the firmware?

- ▶ Your own sketch with TensorFlow Lite Micro.
- ▶ The library of Edge Impulse.

A good Decision Log:

- ▶ gives the flash, the RAM, and the latency of the two paths from your table,
- ▶ says what the latency means for the time that the processor can sleep,
- ▶ names what you lose with your decision. Example: control of the code against the speed of optimized kernels.

If a part does not work, and common problems

Problem	Fix or fallback
Your group has no dataset of Day 2	Copy the folder data/ of a different group into Labs/day02/. Write this in the report.
The inference code is not complete in time	Use the sketch of solutions/ with your three files of the notebook. Continue with Part C.
The upload fails after Day 2	MicroPython is on the board. Start bootloader mode, then upload.
Didn't find op for builtin opcode	The resolver does not name an operator of the model (task B1).
AllocateTensors() failed	The arena is too small. Increase kTensorArenaSize.
The self-test prints FAIL	Check the normalization in task B3. Run section 8 of the notebook again.
The board shows uncertain	The largest probability is below 0.6. Make the motion as in your Day 2 recordings.
The board shows the wrong class for all motions	The model learned from the fallback dataset. Train with your own recordings.
The library of Edge Impulse does not build	Use the estimates of the Studio, and write "estimate".

The README of the lab has the complete list.

- ▶ *Machine Learning Systems* by Vijay Janapa Reddi and contributors, Harvard University. mlsysbook.ai, CC BY-NC-SA 4.0.
- ▶ *TinyML Made Easy: XIAO ESP32S3* by Marcelo Rovai. github.com/Mjrovai/XIAO-ESP32S3-Sense, Apache-2.0.